

kavaskar
2347230

Lab 3

1. Data Preprocessing:

```
In [9]: import numpy as np
import tensorflow as tf
from tensorflow.keras import datasets, utils, layers, models
from tensorflow.keras.preprocessing.image import ImageDataGenerator

# Load CIFAR-10 dataset
(x_train, y_train), (x_test, y_test) = datasets.cifar10.load_data()

# Normalize pixel values to the range [0, 1]
x_train = x_train.astype('float32') / 255.0
x_test = x_test.astype('float32') / 255.0

# One-hot encode class labels
y_train = utils.to_categorical(y_train, 10)
y_test = utils.to_categorical(y_test, 10)

# Data Augmentation (optional)
datagen = ImageDataGenerator(
    horizontal_flip=True,
    rotation_range=15,
    width_shift_range=0.1,
    height_shift_range=0.1
)
datagen.fit(x_train)
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>
170498071/170498071 [=====] - 55s 0us/step

2. Network Architecture Design:

```
In [10]: # Define the feedforward neural network model
model = models.Sequential()

# Input Layer (32x32x3 for CIFAR-10 images)
model.add(layers.InputLayer(input_shape=(32, 32, 3)))

# Hidden Layers
model.add(layers.Conv2D(32, (3, 3), activation='relu', padding='same'))
model.add(layers.Conv2D(32, (3, 3), activation='relu', padding='same'))
model.add(layers.MaxPooling2D(pool_size=(2, 2)))

model.add(layers.Conv2D(64, (3, 3), activation='relu', padding='same'))
model.add(layers.Conv2D(64, (3, 3), activation='relu', padding='same'))
model.add(layers.MaxPooling2D(pool_size=(2, 2)))

model.add(layers.Flatten())
```

```

# Dense Layer before output
model.add(layers.Dense(128, activation='relu'))

# Output Layer with Softmax activation
model.add(layers.Dense(10, activation='softmax'))

# Compile the model
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Summary of the model
model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	896
conv2d (Conv2D)	(None, 32, 32, 32)	896
conv2d_1 (Conv2D)	(None, 32, 32, 32)	9248
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_2 (Conv2D)	(None, 16, 16, 64)	18496
conv2d_3 (Conv2D)	(None, 16, 16, 64)	36928
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 64)	0
flatten (Flatten)	(None, 4096)	0
dense (Dense)	(None, 128)	524416
dense_1 (Dense)	(None, 10)	1290
Total params: 591274 (2.26 MB)		
Trainable params: 591274 (2.26 MB)		
Non-trainable params: 0 (0.00 Byte)		

Justification:

- **Convolutional Layers:** Used for extracting features from images (edges, textures, patterns). They capture spatial hierarchies in images.
- **MaxPooling Layers:** Reduces spatial dimensions, helping prevent overfitting and reducing computational complexity.
- **Dense Layer:** Acts as the final classifier after flattening the feature maps.
- **Softmax Output Layer:** Ideal for multi-class classification (10 output classes).

3. Activation Functions:

We use **ReLU** for hidden layers and **Softmax** for the output layer.

- **ReLU**: Effective in avoiding vanishing gradient issues and accelerating convergence during training. It activates neurons selectively by outputting zero for negative inputs, which helps in sparse representations.
- **Softmax**: Converts raw output scores into probabilities for multi-class classification.

4. Loss Function and Optimizer:

```
In [11]: # Loss function comparison
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accu
```

- **Categorical Cross Entropy**: Suitable for multi-class classification by comparing predicted probabilities with true labels.
- **Adam Optimizer**: Combines the advantages of momentum and RMSprop, adjusting learning rates dynamically. This results in faster and more stable convergence.

Effect of Learning Rate: A high learning rate may lead to unstable training and overshooting the minimum. A low rate can slow down convergence. Adjust the learning rate dynamically if the model does not converge (e.g., using learning rate schedulers).

5. Training the Model:

```
In [12]: # Train the model
history = model.fit(datagen.flow(x_train, y_train, batch_size=64),
                    epochs=50,
                    validation_data=(x_test, y_test))
```

Epoch 1/50
782/782 [=====] - 112s 141ms/step - loss: 1.5246 - accuracy: 0.4413 - val_loss: 1.1981 - val_accuracy: 0.5722

Epoch 2/50
782/782 [=====] - 95s 121ms/step - loss: 1.1290 - accuracy: 0.5975 - val_loss: 1.0090 - val_accuracy: 0.6414

Epoch 3/50
782/782 [=====] - 99s 126ms/step - loss: 0.9787 - accuracy: 0.6569 - val_loss: 0.8212 - val_accuracy: 0.7113

Epoch 4/50
782/782 [=====] - 88s 113ms/step - loss: 0.8905 - accuracy: 0.6872 - val_loss: 0.8078 - val_accuracy: 0.7192

Epoch 5/50
782/782 [=====] - 58s 74ms/step - loss: 0.8221 - accuracy: 0.7129 - val_loss: 0.7506 - val_accuracy: 0.7417

Epoch 6/50
782/782 [=====] - 58s 74ms/step - loss: 0.7778 - accuracy: 0.7277 - val_loss: 0.7919 - val_accuracy: 0.7266

Epoch 7/50
782/782 [=====] - 59s 75ms/step - loss: 0.7412 - accuracy: 0.7403 - val_loss: 0.7785 - val_accuracy: 0.7383

Epoch 8/50
782/782 [=====] - 60s 77ms/step - loss: 0.7040 - accuracy: 0.7506 - val_loss: 0.6992 - val_accuracy: 0.7698

Epoch 9/50
782/782 [=====] - 65s 83ms/step - loss: 0.6824 - accuracy: 0.7625 - val_loss: 0.6806 - val_accuracy: 0.7690

Epoch 10/50
782/782 [=====] - 62s 79ms/step - loss: 0.6575 - accuracy: 0.7683 - val_loss: 0.6780 - val_accuracy: 0.7769

Epoch 11/50
782/782 [=====] - 59s 76ms/step - loss: 0.6380 - accuracy: 0.7776 - val_loss: 0.6859 - val_accuracy: 0.7660

Epoch 12/50
782/782 [=====] - 60s 77ms/step - loss: 0.6261 - accuracy: 0.7804 - val_loss: 0.6662 - val_accuracy: 0.7814

Epoch 13/50
782/782 [=====] - 62s 79ms/step - loss: 0.6077 - accuracy: 0.7860 - val_loss: 0.6281 - val_accuracy: 0.7830

Epoch 14/50
782/782 [=====] - 62s 79ms/step - loss: 0.5984 - accuracy: 0.7905 - val_loss: 0.6572 - val_accuracy: 0.7826

Epoch 15/50
782/782 [=====] - 61s 78ms/step - loss: 0.5802 - accuracy: 0.7961 - val_loss: 0.6583 - val_accuracy: 0.7867

Epoch 16/50
782/782 [=====] - 61s 78ms/step - loss: 0.5745 - accuracy: 0.7989 - val_loss: 0.6501 - val_accuracy: 0.7861

Epoch 17/50
782/782 [=====] - 62s 79ms/step - loss: 0.5592 - accuracy: 0.8030 - val_loss: 0.6500 - val_accuracy: 0.7849

Epoch 18/50
782/782 [=====] - 61s 78ms/step - loss: 0.5521 - accuracy: 0.8081 - val_loss: 0.6049 - val_accuracy: 0.8010

Epoch 19/50
782/782 [=====] - 61s 78ms/step - loss: 0.5365 - accuracy: 0.8119 - val_loss: 0.6358 - val_accuracy: 0.7954

Epoch 20/50
782/782 [=====] - 62s 79ms/step - loss: 0.5359 - accuracy: 0.8134 - val_loss: 0.5642 - val_accuracy: 0.8106

Epoch 21/50
782/782 [=====] - 62s 79ms/step - loss: 0.5254 - accuracy: 0.8159 - val_loss: 0.6118 - val_accuracy: 0.8001

Epoch 22/50
782/782 [=====] - 61s 78ms/step - loss: 0.5170 - accuracy: 0.8193 - val_loss: 0.5803 - val_accuracy: 0.8088

Epoch 23/50
782/782 [=====] - 61s 78ms/step - loss: 0.5119 - accuracy: 0.8202 - val_loss: 0.5747 - val_accuracy: 0.8125

Epoch 24/50
782/782 [=====] - 61s 79ms/step - loss: 0.5090 - accuracy: 0.8241 - val_loss: 0.5896 - val_accuracy: 0.8071

Epoch 25/50
782/782 [=====] - 61s 78ms/step - loss: 0.5051 - accuracy: 0.8239 - val_loss: 0.5744 - val_accuracy: 0.8099

Epoch 26/50
782/782 [=====] - 61s 78ms/step - loss: 0.4967 - accuracy: 0.8251 - val_loss: 0.5748 - val_accuracy: 0.8090

Epoch 27/50
782/782 [=====] - 72s 92ms/step - loss: 0.4881 - accuracy: 0.8293 - val_loss: 0.5895 - val_accuracy: 0.8087

Epoch 28/50
782/782 [=====] - 61s 78ms/step - loss: 0.4823 - accuracy: 0.8315 - val_loss: 0.5843 - val_accuracy: 0.8096

Epoch 29/50
782/782 [=====] - 61s 78ms/step - loss: 0.4755 - accuracy: 0.8343 - val_loss: 0.6243 - val_accuracy: 0.8065

Epoch 30/50
782/782 [=====] - 61s 78ms/step - loss: 0.4768 - accuracy: 0.8317 - val_loss: 0.5667 - val_accuracy: 0.8127

Epoch 31/50
782/782 [=====] - 61s 78ms/step - loss: 0.4709 - accuracy: 0.8339 - val_loss: 0.6211 - val_accuracy: 0.8092

Epoch 32/50
782/782 [=====] - 61s 78ms/step - loss: 0.4663 - accuracy: 0.8366 - val_loss: 0.6184 - val_accuracy: 0.8008

Epoch 33/50
782/782 [=====] - 61s 77ms/step - loss: 0.4637 - accuracy: 0.8365 - val_loss: 0.5548 - val_accuracy: 0.8201

Epoch 34/50
782/782 [=====] - 64s 82ms/step - loss: 0.4564 - accuracy: 0.8384 - val_loss: 0.5650 - val_accuracy: 0.8174

Epoch 35/50
782/782 [=====] - 76s 97ms/step - loss: 0.4520 - accuracy: 0.8418 - val_loss: 0.6222 - val_accuracy: 0.8065

Epoch 36/50
782/782 [=====] - 75s 96ms/step - loss: 0.4507 - accuracy: 0.8415 - val_loss: 0.6510 - val_accuracy: 0.7940

Epoch 37/50
782/782 [=====] - 63s 81ms/step - loss: 0.4470 - accuracy: 0.8427 - val_loss: 0.5701 - val_accuracy: 0.8197

Epoch 38/50
782/782 [=====] - 61s 78ms/step - loss: 0.4411 - accuracy: 0.8445 - val_loss: 0.5847 - val_accuracy: 0.8171

Epoch 39/50
782/782 [=====] - 62s 79ms/step - loss: 0.4373 - accuracy: 0.8490 - val_loss: 0.6096 - val_accuracy: 0.8090

Epoch 40/50
782/782 [=====] - 62s 79ms/step - loss: 0.4445 - accuracy: 0.8425 - val_loss: 0.5912 - val_accuracy: 0.8168

```

Epoch 41/50
782/782 [=====] - 62s 79ms/step - loss: 0.4351 - accuracy: 0.8469 - val_loss: 0.6207 - val_accuracy: 0.8115
Epoch 42/50
782/782 [=====] - 62s 80ms/step - loss: 0.4301 - accuracy: 0.8482 - val_loss: 0.6001 - val_accuracy: 0.8139
Epoch 43/50
782/782 [=====] - 63s 81ms/step - loss: 0.4261 - accuracy: 0.8510 - val_loss: 0.5845 - val_accuracy: 0.8182
Epoch 44/50
782/782 [=====] - 61s 78ms/step - loss: 0.4258 - accuracy: 0.8530 - val_loss: 0.5808 - val_accuracy: 0.8143
Epoch 45/50
782/782 [=====] - 62s 79ms/step - loss: 0.4210 - accuracy: 0.8514 - val_loss: 0.5544 - val_accuracy: 0.8269
Epoch 46/50
782/782 [=====] - 64s 82ms/step - loss: 0.4182 - accuracy: 0.8524 - val_loss: 0.5788 - val_accuracy: 0.8179
Epoch 47/50
782/782 [=====] - 62s 79ms/step - loss: 0.4194 - accuracy: 0.8518 - val_loss: 0.5496 - val_accuracy: 0.8257
Epoch 48/50
782/782 [=====] - 63s 80ms/step - loss: 0.4164 - accuracy: 0.8535 - val_loss: 0.5808 - val_accuracy: 0.8188
Epoch 49/50
782/782 [=====] - 62s 80ms/step - loss: 0.4098 - accuracy: 0.8555 - val_loss: 0.5763 - val_accuracy: 0.8173
Epoch 50/50
782/782 [=====] - 62s 79ms/step - loss: 0.4114 - accuracy: 0.8544 - val_loss: 0.5887 - val_accuracy: 0.8160

```

Backpropagation:

- During backpropagation, weights are updated using gradients from the loss function. The learning rate controls how much to adjust weights. If learning is too slow, increase the rate; if oscillating, decrease it.

6. Model Evaluation:

```

In [13]: # Evaluate the model
test_loss, test_acc = model.evaluate(x_test, y_test)
print(f"Test accuracy: {test_acc}")

# Confusion Matrix and classification report
from sklearn.metrics import confusion_matrix, classification_report
y_pred = model.predict(x_test)
y_pred_classes = np.argmax(y_pred, axis=1)
y_true = np.argmax(y_test, axis=1)

print(confusion_matrix(y_true, y_pred_classes))
print(classification_report(y_true, y_pred_classes))

```

313/313 [=====] - 3s 11ms/step - loss: 0.5887 - accuracy: 0.8160

Test accuracy: 0.8159999847412109

313/313 [=====] - 3s 10ms/step

```
[[847 15 37 3 13 5 5 9 26 40]
 [ 7 928 0 1 1 1 4 1 5 52]
 [ 48 2 737 18 47 40 69 23 4 12]
 [ 20 11 47 562 54 136 90 33 13 34]
 [ 13 2 39 22 781 18 44 71 5 5]
 [ 7 5 31 78 34 749 40 44 2 10]
 [ 6 4 17 8 17 11 915 6 5 11]
 [ 11 5 19 17 27 32 5 868 3 13]
 [ 81 23 4 2 1 1 6 3 849 30]
 [ 20 36 3 2 2 0 3 3 7 924]]
```

	precision	recall	f1-score	support
0	0.80	0.85	0.82	1000
1	0.90	0.93	0.91	1000
2	0.79	0.74	0.76	1000
3	0.79	0.56	0.66	1000
4	0.80	0.78	0.79	1000
5	0.75	0.75	0.75	1000
6	0.77	0.92	0.84	1000
7	0.82	0.87	0.84	1000
8	0.92	0.85	0.88	1000
9	0.82	0.92	0.87	1000
accuracy			0.82	10000
macro avg	0.82	0.82	0.81	10000
weighted avg	0.82	0.82	0.81	10000

- **Precision, Recall, F1-score:** Help identify the model's effectiveness in distinguishing between classes.
- **Confusion Matrix:** Highlights misclassification between classes.

Improving Accuracy: Consider deeper architectures (more layers), data augmentation, or tuning hyperparameters (batch size, learning rate).

7. Optimization Strategies:

- **Early Stopping:** Monitors validation loss to stop training if it stops improving.
- **Learning Rate Scheduling:** Reduces the learning rate gradually to ensure smooth convergence.
- **Weight Initialization:** Proper initialization (e.g., He initialization) avoids vanishing/exploding gradients.

Importance of Weight Initialization: Poor initialization can lead to slow or unstable convergence. Proper initialization helps the network converge faster.

8. Report:

- **Model Architecture:** Detailed explanation and rationale for choices.
- **Training/Test Accuracy:** Relevant plots for loss/accuracy over epochs.

- **Hyperparameters:** Learning rate, batch size, number of epochs, etc.
- **Challenges:** E.g., managing overfitting, tuning hyperparameters.